CS 3368 Introduction to Artificial Intelligence Markov Decision Process

Department of Computer Science Texas Tech University

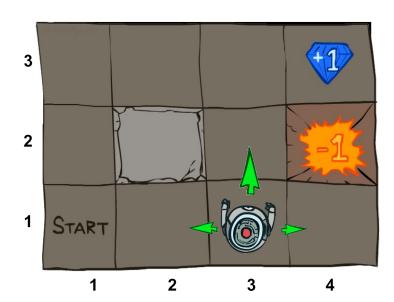
- Instructor: Jingjing Yao
- Email: jingjing.yao@ttu.edu
 - Office: EC 306F
- Office hours: 10 11 am, Tuesday and Thursday



Example: Grid World

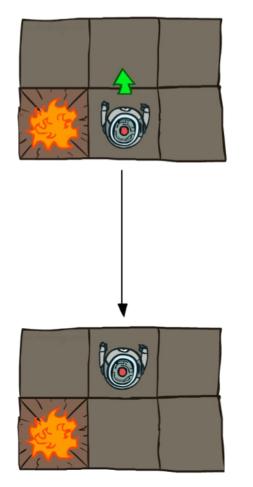
- A maze-like problem
 - The agent lives in a grid
 - Walls block the agent's path
- Noisy movement: actions do not always go as planned
 - 80% of the time, the action North takes the agent North (if there is no wall there)
 - 10% of the time, North takes the agent West; 10% East
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives rewards each time step
 - Small reward each step (can be negative)
 - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards



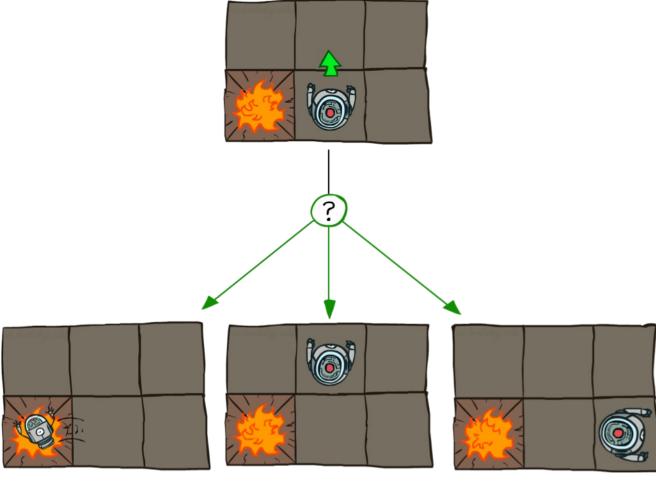


Grid World Actions

Deterministic Grid World



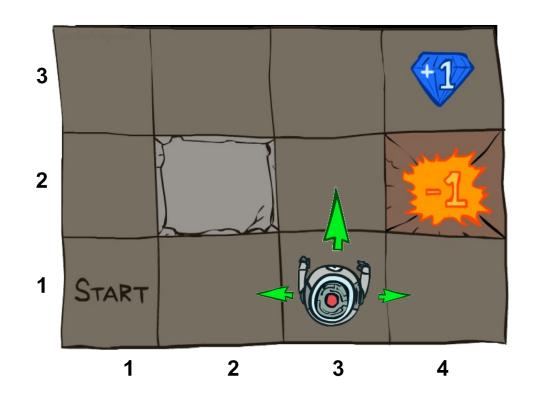
Stochastic Grid World





Markov Decision Processes

- An MDP is defined by:
 - \rightarrow A set of states $s \in S$
 - \rightarrow A set of actions $a \in A$
 - A transition function T(s, a, s')
 - > Probability that a from s leads to s', i.e., P(s'|s,a)
 - > Also called the model or the dynamics
 - A reward function R(s, a, s')
 - \rightarrow Sometimes just R(s) or R(s')
 - > A start state
 - > Maybe a terminal state
- MDPs are non-deterministic search problems
 - > One way to solve them is with expectimax search
 - > We'll have a new tool soon





What is Markov about MDPs?

- "Markov" generally means that given the present state, the future and the past are independent
- For Markov decision processes, "Markov" means action outcomes depend only on the current state

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots S_0 = s_0)$$

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$

 This is just like search, where the successor function could only depend on the current state (not the history)

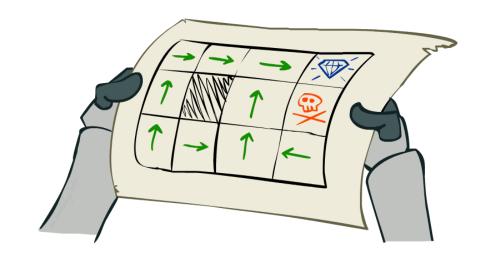


Andrey Markov (1856-1922)



Policies

- In deterministic single-agent search problems, we wanted an optimal plan, or sequence of actions, from start to a goal
- For MDPs, we want an optimal policy $\pi^*: S \to A$
 - > A policy π gives an action for each state
 - An optimal policy is one that maximizes expected utility if followed



Optimal policy when R(s, a, s') = -0.03for all non-terminals s

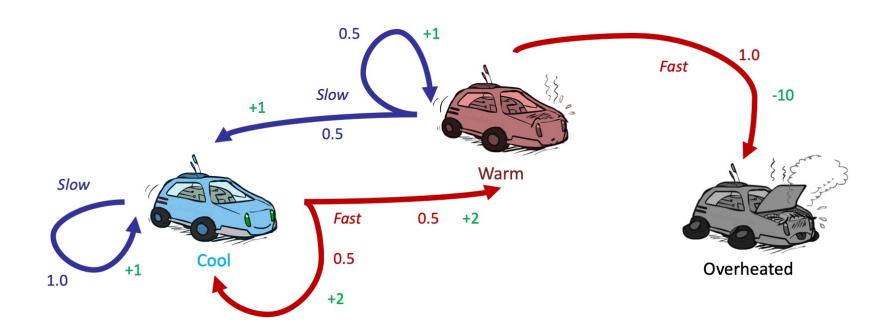


Example: Racing

• A robot car wants to travel far, quickly

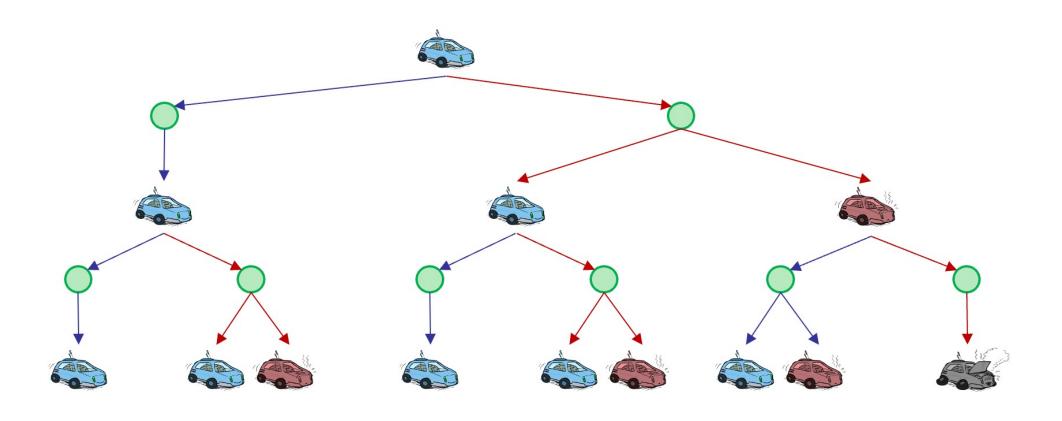
• Three states: Cool, Warm, Overheated

• Two actions: Slow, Fast





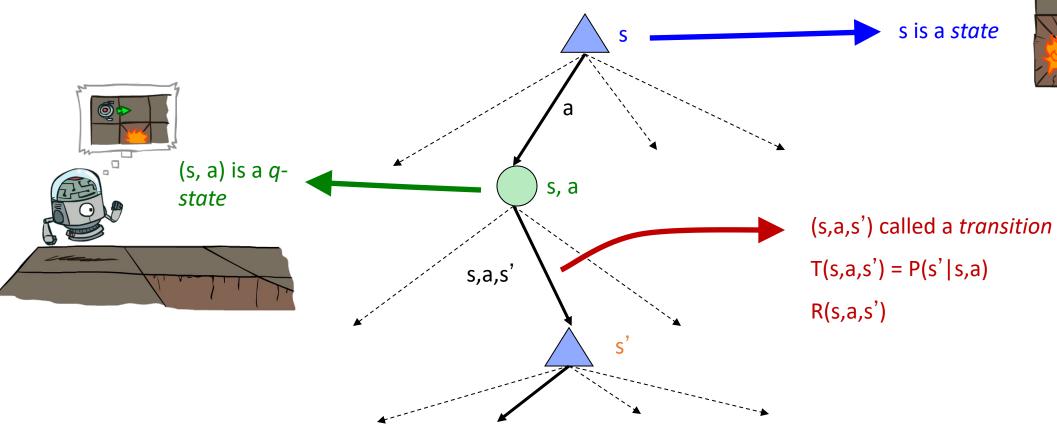
Racing Search Tree





MDP Search Trees

• Each MDP state projects an expectimax-like search tree



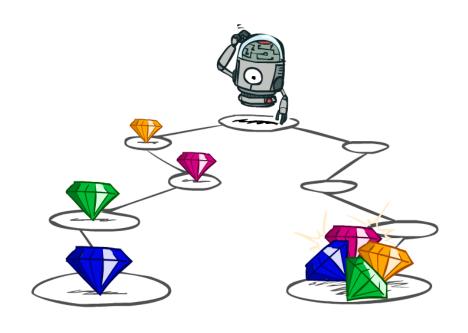


Utilities of Sequences

 What preferences should an agent have over reward sequences?

• More or less? [1, 2, 2] or [2, 3, 4]

• Now or later? [0, 0, 1] or [1, 0, 0]





Discounting

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- One solution: values of rewards decay exponentially

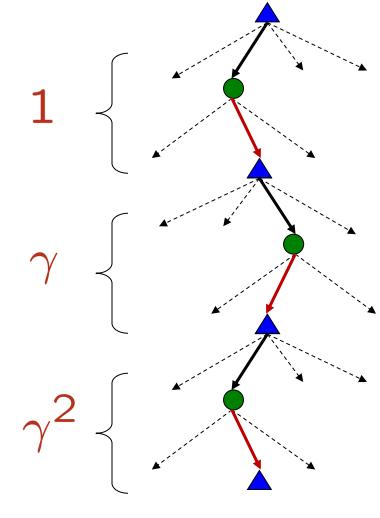




Discounting

- How to discount?
 - Each time we descend a level, we multiply in the discount once
- Why discount?
 - Sooner rewards probably do have higher utility than later rewards
 - Also helps our algorithms converge
- Example: discount of 0.5
 - > U([1,2,3]) = 1*1 + 0.5*2 +
 0.25*3
 - > U([1,2,3]) < U([3,2,1])
- Discounted utility:

$$U([r_0, r_1, r_2, \ldots]) = r_0 + \gamma r_1 + \gamma^2 r_2 \cdots$$

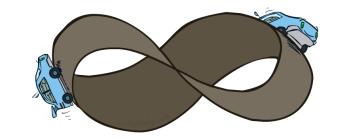




Infinite Utilities?

- Problem: What if the game lasts forever? Do we get infinite rewards?
- Solutions:
 - Finite horizon: Terminate episodes after a fixed T steps (e.g. life)
 - Discounting: use $0 < \gamma < 1$

$$U([r_0, \dots r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \le R_{\text{max}}/(1-\gamma)$$



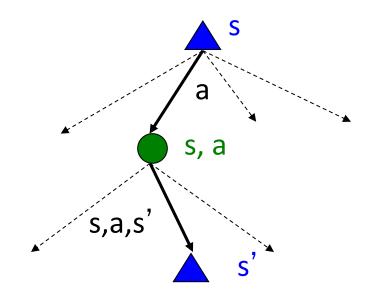
- Smaller γ means smaller "horizon" shorter term focus
- Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like "overheated" for racing)



Summary: Defining MDPs

Markov decision processes:

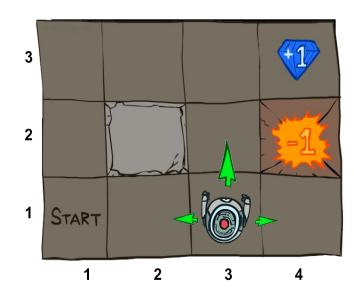
- > Set of states S
- \rightarrow Start state s_0
- > Set of actions A
- > Transitions P(s'|s,a) (or T(s,a,s'))
- > Rewards R(s,a,s') (and discount γ)
- MDP quantities so far:
 - > Policy = Choice of action for each state
 - > Utility = sum of (discounted) rewards





MDP for Grid World

- · Actions: up, down, left, and right
 - > Every action is possible in every state
- The transition model P(s'|s,a)
 - > An action achieves its intended effect with probability 0.8
 - An action leads to a 90-degree left turn with probability 0.1
 - An action leads to a 90-degree right turn with probability 0.1
 - > If the robot bumps into a wall, it stays in the same square
- The reward function R(s) is the reward of entering state s
 - > R(s24) = -1
 - > R(s34) = 1
 - \rightarrow Otherwise, R(s) = -0.04



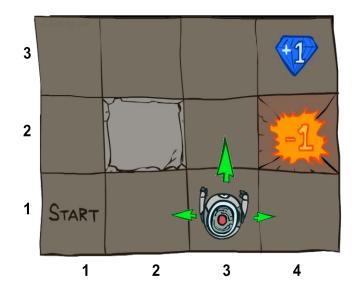


Grid World Problem

What should the robot do to maximize its rewards?

	1	2	3	4
1	Start			
2		X		-1
3				+1

- ▶ Let s_{ij} be the position in row i and column j.
- s₁₁ is the initial state.
- ▶ There is a wall at s₂₂.
- s₂₄ and s₃₄ are goal states.
 The robot escapes the world at either goal state.





Understand Transition Model

CQ: The robot is in s_{14} and tries to move to our right, what is the probability that the robot stays in s_{14} ?

- (A) 0.1
- (B) 0.2
- (C) 0.8
- (D) 0.9
- (E) 1.0

	1	2	3	4
1	Start			
2		X		-1
3				+1



A fixed sequence of actions

CQ: If the environment is deterministic, an optimal solution to the grid world problem is the fixed action sequence: down, down, right, right, and right.

- (A) True
- (B) False
- (C) I don't know

	1	2	3	4
1	Start			
2		X		-1
3				+1



A fixed sequence of actions

CQ: Consider the action sequence "down, down, right, right, and right". This action sequence could take the robot to more than one square with positive probability.

- (A) True
- (B) False
- (C) I don't know

	1	2	3	4
1	Start			
2		Χ		-1
3				+1



The optimal policies of the grid world

 The optimal policy of the grid world changes based on R(s) for any nongoal state s. It shows a careful balancing of risk and reward

	1	2	3	4
1	Start			
2		X		-1
3				+1



The optimal policy

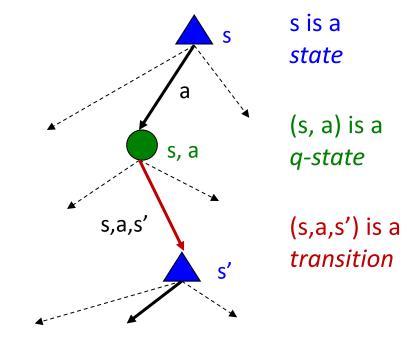
- When the reward function is
- R(s) < -1.6284
- R(s) > -0.02

	1	2	3	4
1	Start			
2		X		-1
3				+1



Optimal Quantities

- The value (utility) of a state s:
 - V*(s) = expected utility starting in s and acting optimally
- The value (utility) of a q-state (s,a):
 Q*(s,a) = expected utility starting out having taken action a from state s and (thereafter) acting optimally
- The optimal policy: $\pi^*(s)$ = optimal action from state s





The Bellman Equations

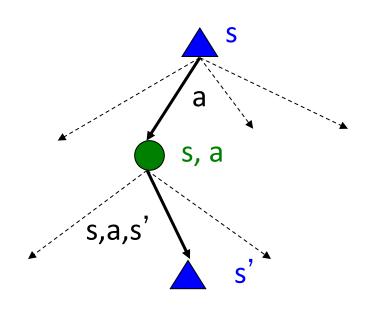
 Definition of "optimal utility" via expectimax recurrence gives a simple one-step lookahead relationship amongst optimal utility values

$$V^{*}(s) = \max_{a} Q^{*}(s, a)$$

$$Q^{*}(s, a) = \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^{*}(s') \right]$$

$$V^{*}(s) = \max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^{*}(s') \right]$$

 These are the Bellman equations, and they characterize optimal values in a way we'll use over and over





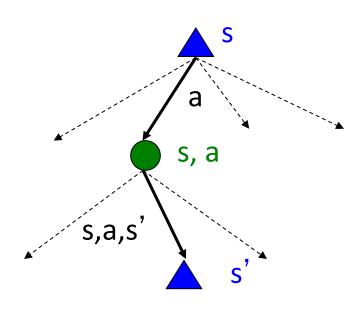
Values of States

- Fundamental operation: compute the (expectimax) value of a state
 - > Expected utility under optimal action
 - > Average sum of (discounted) rewards
 - > This is just what expectimax computed
- Recursive definition of value (Bellman equations)

$$V^{*}(s) = \max_{a} Q^{*}(s, a)$$

$$Q^{*}(s, a) = \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^{*}(s') \right]$$

$$V^{*}(s) = \max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^{*}(s') \right]$$





Calculate the Optimal Policy

What is my expected utility if I am in state s and take action a?

$$Q^{*}(s,a) = \sum_{s'} T(s,a,s') \left[R(s,a,s') + \gamma V^{*}(s') \right]$$

• In state s, choose an action that maximizes my expected utility

$$\pi^*(s) = \arg\max_a Q^*(s, a)$$



Determine optimal action

CQ: What is the optimal action for state s_{13} ?

- (A) Up (B) Down (C) Left (D) Right

$$Q^*(s, a) = \sum_{s'} P(s'|s, a) V^*(s')$$
$$\pi(s) = \arg\max_{a} Q^*(s, a).$$

The values of $V^*(s)$ are given below.

	1	2	3	4
1	0.705	0.655	0.611	0.388
2	0.762	X	0.660	-1
3	0.812	0.868	0.918	+1



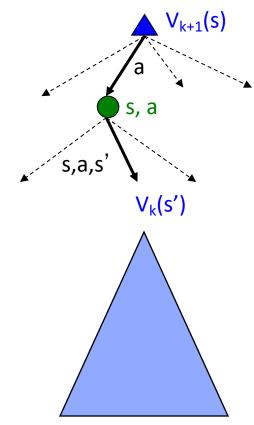
Value Iteration

• Start with $V_0(s) = 0$

• Given vector of $V_k(s)$ values, do one step of expectimax from each state:

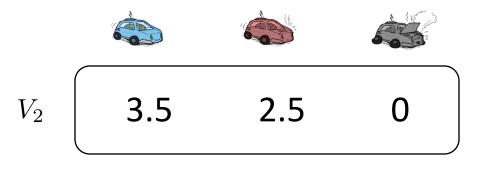
$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V_k(s') \right]$$

- Repeat until convergence
- Theorem: will converge to unique optimal values
 - > Basic idea: approximations get refined towards optimal values
 - > Policy may converge long before values do

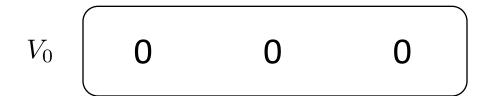


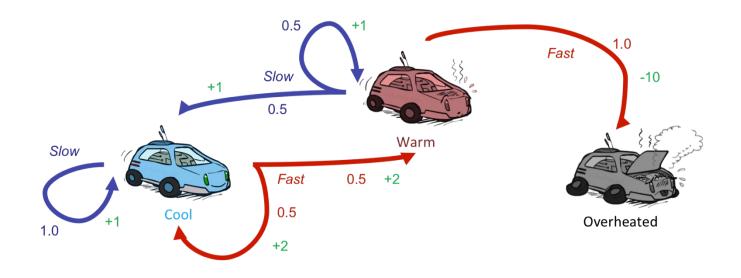


Example: Value Iteration









Assume no discount

$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V_k(s') \right]$$



Example

Write down the Bellman equation for V*(s11)

	1	2	3	4
1	0.705	0.655	0.611	0.388
2	0.762	X	0.660	-1
3	0.812	0.868	0.918	+1



Example

CQ: What is $V_1(s_{23})$?

(A) $(-\infty,0)$ (B) [0,0.25) (C) [0.25,0.5) (D) [0.5,0.75) (E) [0.75,1]

 $V_0(s)$:

	1	2	3	4
1	0	0	0	0
2	0	X	0	-1
3	0	0	0	+1



The Values of $V_1(s)$

 $V_0(s)$:

	1	2	3	4
1	0	0	0	0
2	0	X	0	-1
3	0	0	0	+1

 $V_1(s)$:

	1	2	3	4
1				
2		X		-1
3				+1

	1	2	3	4
1	-0.04	-0.04	-0.04	-0.04
2	-0.04	X	-0.04	-1
3	-0.04	-0.04	0.76	+1



 $V_1(s)$:

	1	2	3	4
1	-0.04	-0.04	-0.04	-0.04
2	-0.04	X	-0.04	-1
3	-0.04	-0.04	0.76	+1

 $V_2(s)$:

	1	2	3	4
1	-0.08	-0.08	-0.08	-0.08
2	-0.08	X	0.464	-1
3	-0.08	0.56	0.832	+1



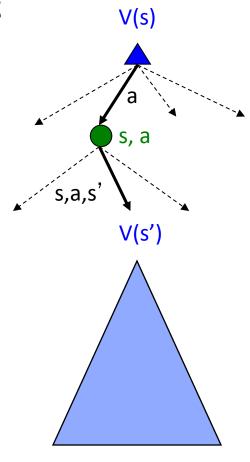
Value Iteration Summary

Bellman equations characterize the optimal values:

$$V^*(s) = \max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^*(s') \right]$$

Value iteration computes them:

$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V_k(s') \right]$$



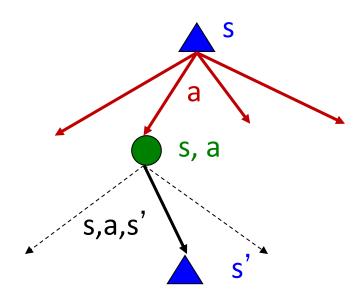


Problems with Value Iteration

Value iteration repeats the Bellman updates:

$$V_{k+1}(s) \leftarrow \max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V_k(s') \right]$$

• Problem 1: It's slow - $O(S^2A)$ per iteration

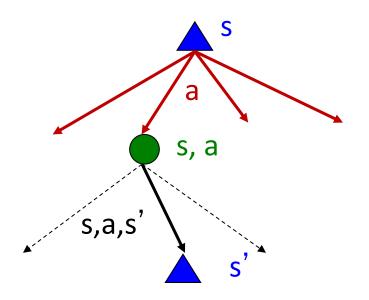


Problem 2: The policy often converges long before the values

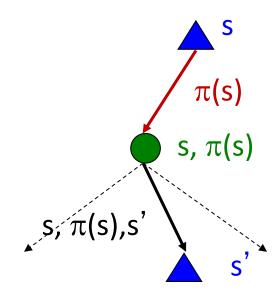


Fixed Policies

Do the optimal action



Do what π says to do



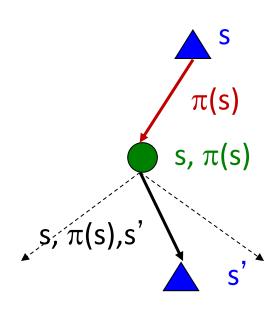
- If we fixed some policy $\pi(s)$, then the tree would be simpler only one action per state
 - > ... though the tree's value would depend on which policy we fixed



Utilities for a Fixed Policy

- Another basic operation: compute the utility of a state s under a fixed (generally non-optimal) policy
- Define the utility of a state s, under a fixed policy π : $V^\pi(s)$ = expected total discounted rewards starting in s and following π
- Recursive relation (one-step look-ahead / Bellman equation):

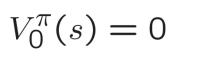
$$V^{\pi}(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^{\pi}(s')]$$



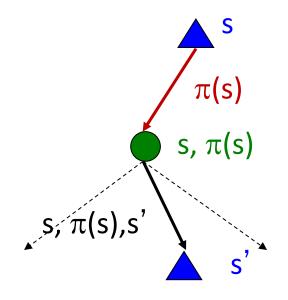


Policy Evaluation

- How do we calculate the V's for a fixed policy π ?
- Idea 1: Turn recursive Bellman equations into updates (like value iteration)



$$V_{k+1}^{\pi}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^{\pi}(s')]$$



- Idea 2: Without the maxes, the Bellman equations are just a linear system
 - > Solve with Matlab (or your favorite linear system solver)



Computing Actions from Values

- Let's imagine we have the optimal values V*(s)
- How should we act?
- We need to do a mini-expectimax (one step)

$$\pi^*(s) = \arg\max_{a} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

 This is called policy extraction, since it gets the policy implied by the values



Computing Actions from Q-Values

Let's imagine we have the optimal Q-values:

How should we act?

$$\pi^*(s) = \arg\max_{a} Q^*(s, a)$$

 Important lesson: actions are easier to select from Q-values than values



Policy Iteration

- Alternative approach for optimal values:
 - Step 1: Policy evaluation: calculate utilities for some fixed policy until convergence
 - Step 2: Policy improvement: update policy using one-step lookahead with resulting converged utilities as future values
 - > Repeat steps until policy converges
- This is policy iteration
 - > It's still optimal
 - > Can converge (much) faster



Policy Iteration

- Evaluation: For fixed current policy π , find values with policy evaluation
 - > Iterate until values converge:

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') \left[R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s') \right]$$

- Improvement: For fixed values, get a better policy using policy extraction
 - > One-step look-ahead:

$$\pi_{i+1}(s) = \arg\max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^{\pi_i}(s') \right]$$



Comparison

- Both value iteration and policy iteration compute the same thing (all optimal values)
- In value iteration:
 - > Every iteration updates both the values and (implicitly) the policy
 - We don't track the policy, but taking the max over actions implicitly recomputes it
- In policy iteration:
 - We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
 - After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
 - > The new policy will be better (or we're done)



Summary: MDP Algorithms

- So you want to....
 - > Compute optimal values: use value iteration or policy iteration
 - > Compute values for a particular policy: use policy evaluation
 - > Turn your values into a policy: use policy extraction
- These all look the same!
 - > They basically are they are all variations of Bellman updates
 - > They all use one-step lookahead expectimax fragments
 - > They differ only in whether we plug in a fixed policy or max over actions



Example

- Reward of entering a non-goal state = -0.04
- Transition probabilities: 0.8 in intended direction, 0.1 to the left and 0.1 to the right
- Execute Policy Iteration
- Initial policy $\pi(s11)$ = right; $\pi(s21)$ = right

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') \left[R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s') \right]$$

$$\pi_{i+1}(s) = \arg\max_{a} \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^{\pi_i}(s') \right]$$

